

data_clustering_1

June 26, 2018

1 Data Clustering

Data clustering is a process of assigning a set of records into subsets, called clusters, such that records in the same cluster are similar and records in different clusters are quite distinct.

A typical clustering process involves the following five steps:

1. pattern representation;
2. dissimilarity measure definition;
3. clustering;
4. data abstraction;
5. assesment of output

In this interactive session, we will be reviewing the *k-means* algorithm.

1.1 The *k-means* Algorithm

The *k-means* algorithm is the most popular and the simplest partitional clustering algorithm. It has many variations. This exercise will review the standard algorithm.

1.1.1 Description of the Algorithm

Let $X = \{\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_{n-1}\}$ be a numeric dataset containing n records and k be an integer in $\{1, 2, \dots, n\}$. The *k-means* algorithm tries to divide the dataset into k clusters C_0, C_1, \dots , and C_{k-1} by minimizing the following objective function:

$$E = \sum_{i=0}^{k-1} \sum_{\mathbf{x} \in C_i} D(\mathbf{x}, \boldsymbol{\mu}_i),$$

where $D(\cdot, \cdot)$ is a distance measure and $\boldsymbol{\mu}_i$ is the mean of cluster C_i , i.e.,

$$\boldsymbol{\mu}_i = \frac{1}{|C_i|} \sum_{\mathbf{x} \in C_i} \mathbf{x}$$

Let γ_i be the cluster membership of record \mathbf{x}_i for $i = 0, 1, \dots, n-1$. That is, $\gamma_i = j$ if \mathbf{x}_i belongs to cluster C_j . Then the objective function can be rewritten as:

$$E = \sum_{i=0}^{n-1} D(\mathbf{x}_i, \boldsymbol{\mu}_{\gamma_i})$$

To minimize the objective function, the *k-means* algorithm employs an iterative process. At the beginning, the *k-means* algorithm selects k random records from the dataset X as initial cluster centers.

Suppose $\mu_0^{(0)}, \mu_1^{(0)}, \dots$, and $\mu_{k-1}^{(0)}$ are the initial cluster centers. Based on these cluster centers, the *k-means* algorithm updates the cluster memberships $\gamma_0^{(0)}, \gamma_1^{(0)}, \gamma_{n-1}^{(0)}$ as follows:

$$\gamma_i^{(0)} = \underset{0 \leq j \leq k-1}{\operatorname{argmin}} D(\mathbf{x}_i, \mu_j^{(0)}) \quad (1)$$

where argmin is the argument that minimizes the distance. That is, $\gamma_i^{(0)}$ is set to the index of the cluster to which \mathbf{x}_i has the smallest distance.

Based on the cluster memberships $\gamma_0^{(0)}, \gamma_1^{(0)}, \gamma_{n-1}^{(0)}$, the *k-means* algorithm updates the cluster centers as follows:

$$\mu_j^{(1)} = \frac{1}{\left| \left\{ i \mid \gamma_i^{(0)} = j \right\} \right|} \sum_{i=0, \gamma_i^{(0)}=j} \mathbf{x}_i, \quad j = 0, 1, \dots, k-1 \quad (2)$$

Then the *k-means* algorithm repeats updating the cluster memberships based on Equation 1 and 2.

1.2 Exercises

1.2.1 To practice your Python

1. From the datasets card in trello, use `data.tar.gz` to get the file with name `600points.csv`.
2. Use the `kmeans` implementation from `scikit-learn` to run the algorithm for $k = 3$ clusters and a max of 100 iterations. You can install that using `pip ($ pip install scikit-learn)` or `anaconda, ($ conda install scikit-learn)`.
3. Use python's `docopt` or `argparse` (available in the standard library) to parse the following command line options:

Allowed options:

```
--help           produce help message
--datafile arg   the data file
--k arg (=3)     number of clusters
--maxiter arg (=100) maximum number of iterations
```

4. Implement the algorithm by yourself and compare with the results obtained with `scikit-learn`.
5. Use `numba` to try to improve the performance of your implementation.

1.2.2 To practice with .NET (C# or F#)

1. Use the *k-means* implementation from version 0.2 of [ML.NET](#) as a reference implementation for comparison (correctness of results, performance, etc).

2. Use [.NET's docopt port](#) to parse command line options for your program (see analogous section for Python).
3. Implement your own version of k-means using your preferred .NET language (C#, F#, VB.NET)

1.2.3 To practice with GPUs and Python/R/C

1. Take a look at [kmcuda](#). See what the possibilities are for your favorite programming language.

1.2.4 To practice with GPUs or Multicore CPUs and Haskell

1. Take a look at [Accelerate](#) and the different supported backends (e.g. Multicore CPUs and GPUs).
2. Study and explain the [k-means sample](#)
3. Compare the implementations and performance.

1.2.5 To practice with Scala and Spark

1. Try to use the existing Scala implementations:
 - With [Spark's MLLib](#)
 - With [Spire](#)
2. Study and explain.
3. Hard: study if it is possible to improve performance using [sbt-javacpp](#) bindings for CUDA and previous implementation code/algorithms.